

# Relational Algebra and SQL

# Relational Query Languages

- Languages for describing queries on a relational database
- *Structured Query Language* (SQL)
  - Predominant application-level query language
  - Declarative
- *Relational Algebra*
  - Intermediate language used within DBMS
  - Procedural

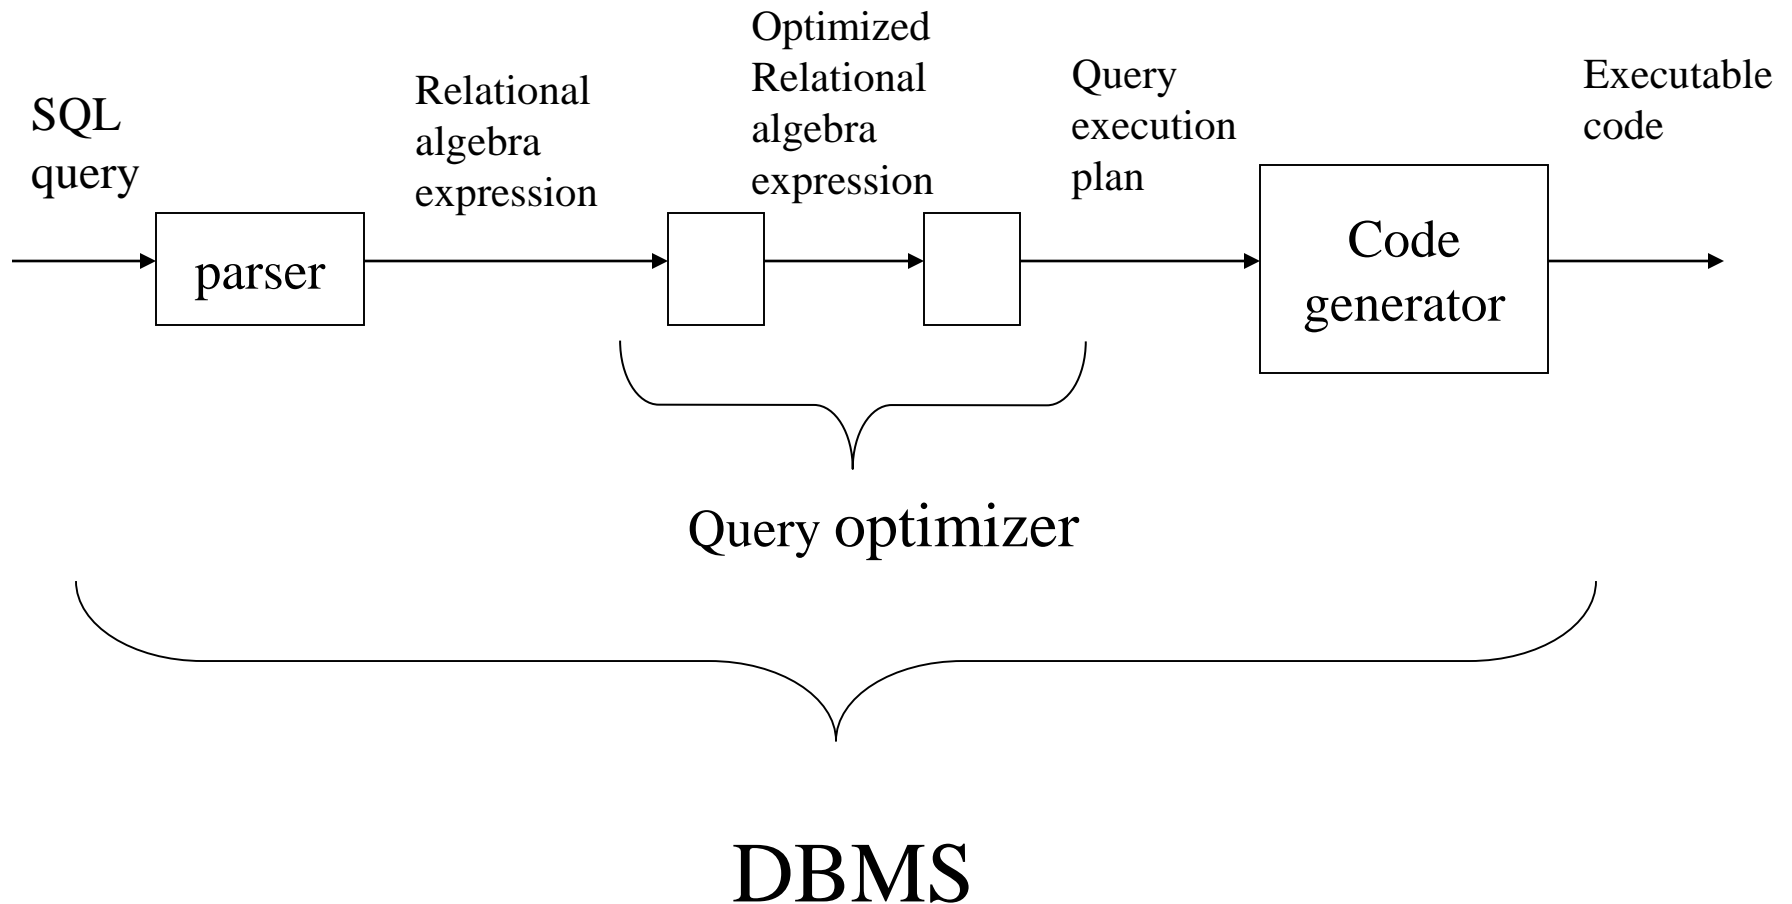
# What is an Algebra?

- A language based on operators and a domain of values
- Operators map values taken from the domain into other domain values
- Hence, an expression involving operators and arguments produces a value in the domain
- When the domain is a set of all relations (and the operators are as described later), we get the *relational algebra*
- We refer to the expression as a *query* and the value produced as the *query result*

# Relational Algebra

- *Domain*: set of relations
- *Basic operators*: select, project, union, set difference, Cartesian product
- *Derived operators*: set intersection, division, join
- *Procedural*: Relational expression specifies query by describing an algorithm (the sequence in which operators are applied) for determining the result of an expression

# Relational Algebra in a DBMS



# Schema for Student Registration System

Student (*Id, Name, Addr, Status*)

Professor (*Id, Name, DeptId*)

Course (*DeptId, CrsCode, CrsName, Descr*)

Transcript (*StudId, CrsCode, Semester, Grade*)

Teaching (*ProfId, CrsCode, Semester*)

Department (*DeptId, Name*)

# Relational Algebra

- Five operators:
  - Selection:  $\sigma$
  - Projection:  $\Pi$
  - Cartesian Product:  $\times$
  - Union:  $\cup$
  - Difference:  $-$
- Derived or auxiliary operators:
  - Intersection, complement
  - Joins (natural, equi-join, theta join, semi-join)
  - Renaming:  $\rho$

# Set Operators

- Relation is a set of tuples, so set operations should apply:  $\cap$ ,  $\cup$ ,  $-$  (set difference)
- Result of combining two relations with a set operator is a relation  $\Rightarrow$  all its elements must be tuples having same structure
- Hence, scope of set operations limited to *union compatible relations*



# Select Operator

- Produce table containing subset of rows of argument table satisfying condition

$\sigma_{condition}$  *relation*

- Example:

Person

<i>Id</i>	<i>Name</i>	<i>Address</i>	<i>Hobby</i>
1123	John	123 Main	stamps
1123	John	123 Main	coins
5556	Mary	7 Lake Dr	hiking
9876	Bart	5 Pine St	stamps

$\sigma_{Hobby='stamps'}(Person)$

<i>Id</i>	<i>Name</i>	<i>Address</i>	<i>Hobby</i>
1123	John	123 Main	stamps
9876	Bart	5 Pine St	stamps

# 1. Select ( $\sigma$ )

- Returns all tuples which satisfy a condition

- Notation:  $\sigma_c(R)$

SSN	Name	Salary
1234545	John	200000
5423341	Smith	600000
4352342	Fred	500000

- Examples

- $\sigma_{\text{Salary} > 40000}(\text{Employee})$

SSN	Name	Salary
5423341	Smith	600000
4352342	Fred	500000

- $\sigma_{\text{name} = \text{"Smith"}}(\text{Employee})$

SSN	Name	Salary
5423341	Smith	600000

The condition c can be =, <, ≤, >, ≥, <>

# Selection Condition

- Operators:  $<$ ,  $\leq$ ,  $\geq$ ,  $>$ ,  $=$ ,  $\neq$
- Conditions:
  - *$\langle \text{attribute} \rangle$  operator  $\langle \text{constant} \rangle$*
  - *$\langle \text{attribute} \rangle$  operator  $\langle \text{attribute} \rangle$*
  - *$\langle \text{condition} \rangle$  AND  $\langle \text{condition} \rangle$*
  - *$\langle \text{condition} \rangle$  OR  $\langle \text{condition} \rangle$*
  - *NOT  $\langle \text{condition} \rangle$*

# Select - Examples

- $\sigma_{Id > 3000 \text{ Or } Hobby = 'hiking'}(Person)$
- $\sigma_{Id > 3000 \text{ AND } Id < 3999}(Person)$
- $\sigma_{NOT(Hobby = 'hiking')}(Person)$
- $\sigma_{Hobby \neq 'hiking'}(Person)$

*Person*

<i>Id</i>	<i>Name</i>	<i>Address</i>	<i>Hobby</i>
1123	John	123 Main	stamps
1123	John	123 Main	coins
5556	Mary	7 Lake Dr	hiking
9876	Bart	5 Pine St	stamps

## 2. Project ( $\Pi$ )

Produces table containing subset of columns of argument table

$$\Pi_{\text{attribute list}}(\text{relation})$$

*Example 1:*

**Person**

<i>Id</i>	<i>Name</i>	<i>Address</i>	<i>Hobby</i>
1123	John	123 Main	stamps
1123	John	123 Main	coins
5556	Mary	7 Lake Dr	hiking
9876	Bart	5 Pine St	stamps

**$\Pi_{\text{Name,Hobby}}(\text{Person})$**

<i>Name</i>	<i>Hobby</i>
John	stamps
John	coins
Mary	hiking
Bart	stamps

# Project Operator

- *Example 2:*

**Person**

<i>Id</i>	<i>Name</i>	<i>Address</i>	<i>Hobby</i>
1123	John	123 Main	stamps
1123	John	123 Main	coins
5556	Mary	7 Lake Dr	hiking
9876	Bart	5 Pine St	stamps

$\Pi_{Name,Address}(\text{Person})$

<i>Name</i>	<i>Address</i>
John	123 Main
Mary	7 Lake Dr
Bart	5 Pine St

Result is a table (no duplicates)

# Expressions

$\Pi_{Id, Name} ( \sigma_{Hobby='stamps' \text{ OR } Hobby='coins'} (Person) )$

<i>Id</i>	<i>Name</i>	<i>Address</i>	<i>Hobby</i>
1123	John	123 Main	stamps
1123	John	123 Main	coins
5556	Mary	7 Lake Dr	hiking
9876	Bart	5 Pine St	stamps

*Person*

<i>Id</i>	<i>Name</i>
1123	John
9876	Bart

*Result*

### 3. Union Operation (U)

- Fetch data from two relations(tables) or temporary relation(result of another operation).
- Relations(tables) specified should have same number of attributes(columns) and compatible.
- Duplicate tuples are eliminated from the result.

**Syntax:**  $R1 \cup R2$

where R1 and R2 are relations.

$\prod_{\text{StdID,Name,StdCourse}}(R1) \cup \prod_{\text{StdID,Name,StdCourse}}(R2)$

StdID	Name	StdCourse
4352342	John	EC201
4352500	Smita	CS200
5423341	Krishna	CS101

StdID	Name	StdCourse
5423341	Krishna	CS101
4352342	John	EC201

R1

StdID	Name	Stdcourse
5423341	Krishna	CS101
4352500	Smita	CS200

R2



# 4. Difference Operation (-)

- Attribute name in relation R1 has to match with attribute name in R2.
- Both Relations should be compatible
- Resultant relation will have tuples present in one relation and not present in the second relation.

**Syntax:**  $R1 - R2$

where R1 and R2 are relations.

$\Pi_{\text{Stdid,Name,StdCourse}}(R1) - \Pi_{\text{Stdid,Name,Stdcourse}}(R2)$

StdID	Name	StdCourse
5423341	Krishna	CS101
4352342	John	EC201

$R1$

StdID	Name	StdCourse
4352342	John	EC201

StdID	Name	Stdcourse
5423341	Krishna	CS101
4352500	Smita	CS200

$R2$

# 4. Intersection Operation ( $\cap$ )

- Attribute name in relation R1 has to match with attribute name in R2.
- Both Relations should be compatible
- Resultant relation will have tuples present in both relations
- 

**Syntax:**  $R1 \cap R2$

where R1 and R2 are relations.

$\Pi_{\text{Stdid, Name, StdCourse}}(R1) \cap \Pi_{\text{Stdid, Name, StdCourse}}(R2)$

StdID	Name	StdCourse
5423341	Krishna	CS101

StdID	Name	StdCourse
5423341	Krishna	CS101
4352342	John	EC201

$R1$

StdID	Name	Stdcourse
5423341	Krishna	CS101
4352500	Smita	CS200

$R2$

# 5. Cartesian Product (X)

- If  $R1$  and  $R2$  are two relations,  $R1 \times R2$  is the set of all concatenated tuples  $\langle x, y \rangle$ , where  $x$  is a tuple in  $R1$  and  $y$  is a tuple in  $R2$ 
  - ( $R1$  and  $R2$  need not be union compatible)
- $R1 \times R2$  is expensive to compute:

<i>a</i>	<i>b</i>	<i>c</i>	<i>d</i>
x1	x2	y1	y2
x3	x4	y3	y4
<b><i>R1</i></b>		<b><i>R2</i></b>	

<i>a</i>	<i>b</i>	<i>c</i>	<i>d</i>
x1	x2	y1	y2
x1	x2	y3	y4
x3	x4	y1	y2
x3	x4	y3	y4

***$R1 \times R2$***

# Cartesian Product

- Concatenation of every tuple of one relation with every tuple of a second relations.
- Relation A (having  $m$  tuples) and relation B (having  $n$  tuples) has  $m$  times  $n$  tuples.
- Denoted  $A \times B$  or  $A \text{ TIMES } B$ .

Notation:  $R1 \times R2$

Example: Student  $\times$  Credit

*Result :  $R1 \times R2$*

*R1*

StdID	Name	Stdcourse
5423341	Krishna	CS101
4352500	Smita	CS200

*R2*

Id	Credit	Dep
5423341	4	CS
4352300	3	CE

StdID	Name	Stdcourse	Id	Credit	Dep
5423341	Krishna	CS101	5423341	4	CS
5423341	Krishna	CS101	4352300	3	CE
4352500	Smita	CS200	5423341	4	CS
4352500	Smita	CS200	4352300	3	CE

$\sigma_{Credit = '4'}(R1 \times R2)$

## 6. Rename ( $\rho$ )

Changes the schema, not the instance

Notation:  $\rho_{B_1, \dots, B_n}(R)$

Example:

Old Schema:

*Student(StdID, Name, StdCourse)*

—  $\rho_{id, StudentName}(Student)$

New schema:

*Student(Id, StudentName)*

StdID	Name	Stdcourse
5423341	Krishna	CS101
4352500	Smita	CS200

id	StudentName	Stdcourse
5423341	Krishna	CS101
4352500	Smita	CS200

# Rename ( $\rho$ )

- Result of expression evaluation is a relation
- Attributes of relation must have distinct names. This is not guaranteed with Cartesian product
  - e.g., suppose in previous example  $a$  and  $c$  have the same name
- Renaming operator tidies this up. To assign the names  $A_1, A_2, \dots, A_n$  to the attributes of the  $n$  column relation produced by expression  $expr$  use  
$$expr [A_1, A_2, \dots, A_n]$$

# Join and Rename

- **Problem:**  $R1$  and  $R2$  might have attributes with the same name – in which case the Cartesian product is not Possible
- **Solution:**
  - Rename attributes prior to forming the product and use new names in *join-condition*.
  - Common attribute names are qualified with relation names in the result of the join

# Relational Algebra and SQL

**Which of the two expression is more easy to understand?**

$\pi_{CrsName} \sigma_{C\_CrsCode=T\_CrsCode \text{ AND } Sem='S2000'}$   
 $(Course [C\_CrsCode, DeptId, CrsName, Desc]$   
 $\times Teaching [ProfId, T\_CrsCode, Sem])$

**equivalent to:**

```
SELECT C.CrsName
FROM Course C, Teaching T
WHERE C.CrsCode=T.CrsCode AND T.Sem='S2000'
```

SELECT is simple evaluation algorithm .



# Derived Operation: Join

The expression :

$$\sigma_{\text{join-condition}} (R1 \times R2)$$

where *join-condition* is a *conjunction* of terms:

$A_i \text{ oper } B$

$A_i$  is an attribute of  $R1$ ,  $B_i$  is an attribute of  $R2$ , and *oper* is one of  $=, <, >, \geq, \neq, \leq$ ,  
is referred to as the (theta) join of  $R1$  and  $R2$  and denoted:

$R1 \text{ join-condition } R2$

where *join-condition* is represented by  $\bowtie$  or  $|X|$

# Theta Join

A join that involves a predicate

$$R1 \bowtie_{\theta} R2 = \sigma_{\theta} (R1 \times R2)$$

- Here  $\theta$  can be any condition

# Theta Join

A (*general* or *theta*) *join* of  $R1$  and  $R2$  is the expression

$$R1 \bowtie_c R2$$

where *join-condition*  $c$  is a *conjunction* of terms:

$$A_i \text{ oper } B_i$$

Where  $A_i$  is an attribute of  $R1$ ;  $B_i$  is an attribute of  $R2$  ; and *oper* is one of  $=, <, >, \geq, \neq, \leq$ .

**Q:** Any difference between join condition and selection condition?

The meaning is:

$$\sigma_c(R1 \times R2)$$

Where join-condition  $c$  becomes a select condition  $c$  except for possible renamings of attributes

# Theta Join – Example

Output the names of all employees that earn more than their managers.

$\Pi_{\text{Employee.Name}} (\text{Employee} \bowtie_{\text{MngrId=Id AND Salary} > \text{Salary}} \text{Manager})$

The join yields a table with attributes:

*Employee.Name*, *Employee.Id*, *Employee.Salary*, *Employee.MngId*,  
*Manager.Name*, *Manager.Id*, *Manager.Salary*

# Natural Join

- Special case of equijoin:
  - join condition equates *all* and *only* those attributes with the same name (condition doesn't have to be explicitly stated)
  - duplicate columns eliminated from the result

Transcript (*StudId*, *CrsCode*, *Sem*, *Grade*)

Teaching (*ProfId*, *CrsCode*, *Sem*)

Transcript  $\bowtie$  Teaching =

$\pi_{StudId, Transcript.CrsCode, Transcript.Sem, Grade, ProfId} ($   
Transcript  $\bowtie_{CrsCode=CrsCode \text{ AND } Sem=Sem}$  Teaching) $)$

# Natural Join (con't)

Notation :

$$R1 \bowtie R2 = \pi_{attr-list} ( \sigma_{join-cond} (R1 \times R2) )$$

where

$attr-list = attributes(R1) \cup attributes(R2)$   
(duplicates are eliminated) and  $join-cond$  has the form:

$$A_1 = A_1 \text{ AND } \dots \text{ AND } A_n = A_n$$

where

$$\{A_1 \dots A_n\} = attributes(R1) \cap attributes(R2)$$

# Natural Join

Notation:  $R1 \bowtie R2$

Meaning:  $R1 \bowtie R2 = \Pi_A(\sigma_C(R1 \times R2))$

Where:

- The selection  $\sigma_C$  checks equality of all common attributes
- The projection eliminates the duplicate common attributes

# Natural Join

- $R1 =$ 

A	B
X	Y
X	Z
Y	Z
Z	V

 $R2 =$ 

B	C
Z	U
V	W
Z	V

- $R1 \bowtie R2 =$ 

A	B	C
X	Z	U
X	Z	V
Y	Z	U
Y	Z	V
Z	V	W



# Natural Join : Examples

- Given the schemas  $R1(A, B, C, D)$ ,  $R2(A, C, E)$ , what is the schema of  $R \bowtie S$  ?
- Given  $R1(A, B, C)$ ,  $R2(D, E)$ , what is  $R1 \bowtie R2$  ?
- Given  $R1(A, B)$ ,  $R2(A, B)$ , what is  $R1 \bowtie R2$  ?